

# Question 1

Question1. (A)

We first get an insight into the data by looking at our base model including all the features, looking at the correlation matrix, OOS  $R^2$ , and identifying which features are significant, highly correlated, etc.

```
library(tidyverse)
library(dplyr)

train <- read.csv("laptop_train.csv")
test  <- read.csv("laptop_test.csv")

train$Company <- factor(train$Company)
train$TypeName <- factor(train$TypeName)
train$GPU <- factor(train$GPU)

# These should stay numeric
num_vars <- c("Screen", "Memory", "Weight", "Rating", "Price")
train[num_vars] <- lapply(train[num_vars], as.numeric)

# correlation matrix
cor(train[, num_vars], use = "complete.obs")
```

	Screen	Memory	Weight	Rating	Price
Screen	1.00000000	0.09940796	0.81150314	-0.03097351	-0.1306660
Memory	0.09940796	1.00000000	0.29239722	0.02431015	0.7224164
Weight	0.81150314	0.29239722	1.00000000	-0.01924711	0.1167789
Rating	-0.03097351	0.02431015	-0.01924711	1.00000000	-0.0290045
Price	-0.13066597	0.72241635	0.11677891	-0.02900450	1.0000000

```
# Base Model
modell1 <- lm(Price ~ InventoryID + Screen + Memory + Weight + Rating + Company + TypeName + GPU, data =
summary(modell1)
```

```
##
## Call:
## lm(formula = Price ~ InventoryID + Screen + Memory + Weight +
##     Rating + Company + TypeName + GPU, data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -975.88 -188.34  -35.98  161.98 1523.80
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   1421.95846   289.06648    4.919 1.10e-06 ***
## InventoryID      0.06703    0.08057    0.832  0.40572
## Screen       -120.38632    21.81619   -5.518 4.94e-08 ***
## Memory         87.25133     4.40165   19.822 < 2e-16 ***
```

```
## Weight          270.53619   49.45684    5.470 6.41e-08 ***
## Rating          -11.66999    4.61720   -2.528 0.01172 *
## CompanyDell      77.66116   44.30153    1.753 0.08007 .
## CompanyHP       147.04460   51.77867    2.840 0.00465 **
## CompanyLenovo    -1.75199   61.33453   -0.029 0.97722
## TypeNameNotebook -81.93259   59.12440   -1.386 0.16629
## TypeNameUltrabook 405.48355   76.00576    5.335 1.32e-07 ***
## GPUIntel        164.06650   39.70629    4.132 4.06e-05 ***
## GPUNvidia       217.76478   46.46989    4.686 3.39e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 337.7 on 652 degrees of freedom
## Multiple R-squared:  0.6601, Adjusted R-squared:  0.6538
## F-statistic: 105.5 on 12 and 652 DF,  p-value: < 2.2e-16
```

```
library(olsrr)
ols_step_backward_p(model1, p_val = 0.05, progress = TRUE)
```

```
## Backward Elimination Method
## -----
```

```
##
## Candidate Terms:
##
```

```
## 1. InventoryID
## 2. Screen
## 3. Memory
## 4. Weight
## 5. Rating
## 6. Company
## 7. TypeName
## 8. GPU
```

```
##
##
## Variables Removed:
```

```
##
## => InventoryID
##
```

```
## No more variables to be removed.
```

```
##
##
```

```
## Stepwise Summary
```

```
## -----
## Step   Variable      AIC      SBC      SBIC      R2      Adj. R2
## -----
## 0      Full Model    9645.396  9708.393  7750.566  0.66009  0.65383
## 1      InventoryID  9644.102  9702.599  7749.215  0.65973  0.65399
## -----
```

```
##
## Final Model Output
## -----
```

```
##
## Model Summary
```

```
## -----
```

```
## R                0.812      RMSE                334.527
## R-Squared        0.660      MSE                  111908.165
## Adj. R-Squared   0.654      Coef. Var            32.882
## Pred R-Squared   0.644      AIC                  9644.102
## MAE              247.405     SBC                  9702.599
```

```
## -----
## RMSE: Root Mean Square Error
## MSE: Mean Square Error
## MAE: Mean Absolute Error
## AIC: Akaike Information Criteria
## SBC: Schwarz Bayesian Criteria
##
```

```
## ANOVA
```

```
## -----
##              Sum of
##              Squares      DF      Mean Square      F      Sig.
## -----
## Regression    144284374.791      11    13116761.345    115.095    0.0000
## Residual      74418929.460     653     113964.670
## Total         218703304.251     664
```

```
## Parameter Estimates
```

```
## -----
##              model      Beta      Std. Error      Std. Beta      t      Sig.      lower      upper
## -----
##      (Intercept)    1434.247      288.621              4.969    0.000      867.510    2000.984
##      Screen        -120.932      21.801      -0.249    -5.547    0.000     -163.741     -78.123
##      Memory         87.373       4.398       0.574    19.866    0.000      78.736     96.009
##      Weight        271.099      49.441       0.287     5.483    0.000     174.017    368.180
##      Rating        -11.821       4.613      -0.059    -2.563    0.011     -20.878     -2.763
##      CompanyDell     86.725      42.931       0.069     2.020    0.044      2.426    171.025
##      CompanyHP      170.395      43.502       0.131     3.917    0.000      84.974    255.816
##      CompanyLenovo   35.328      42.129       0.028     0.839    0.402     -47.396    118.052
##      TypeNameNotebook -76.295      58.721     -0.061    -1.299    0.194    -191.600     39.009
##      TypeNameUltrabook 416.394      74.848       0.265     5.563    0.000     269.421    563.366
##      GPUIntel       165.430      39.663       0.144     4.171    0.000      87.547    243.312
##      GPUNvidia      218.139      46.457       0.173     4.696    0.000     126.916    309.361
## -----
```

```
# out-of-sample R2 for base model
pred_test <- predict(model1, newdata = test)
sse <- sum((test$Price - pred_test)^2)
sst <- sum((test$Price - mean(test$Price))^2)
R2_out <- 1 - sse/sst
R2_out
```

```
## [1] 0.5506981
```

Looking at the correlation matrix, I observed that Screen and Weight are highly correlated (0.81). This high collinearity inflates variances of coefficient estimates and makes interpretation unstable. Indeed, in the full regression output, both Screen and Weight appeared significant, but their strong correlation makes it difficult to separate their individual contributions. To avoid redundancy and multicollinearity, I decided to drop Weight and retain Screen, which has a closer correlation to our dependent variable (price), low p-value (statistically significant), and which is more directly

interpretable as a feature consumers see when buying laptops. Next, I examined the coefficient for InventoryID. This variable is simply an internal stock identifier and has no managerial meaning for pricing. Its coefficient was very small (0.067), with a p-value of 0.406, confirming it was not statistically significant. Including InventoryID risks overfitting without adding explanatory value. Therefore, I chose to exclude InventoryID from the second model.

```
# 2nd Model to compare with
model2 <- lm(Price ~ Screen + Memory + Rating + Company + TypeName + GPU, data = train)
summary(model2)

##
## Call:
## lm(formula = Price ~ Screen + Memory + Rating + Company + TypeName +
##     GPU, data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1040.73  -194.38   -36.05   166.39  1559.87
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      823.027    272.079   3.025  0.00258 **
## Screen          -36.598     15.791  -2.318  0.02077 *
## Memory           90.206      4.464  20.208 < 2e-16 ***
## Rating          -12.815      4.710  -2.721  0.00669 **
## CompanyDell      124.905     43.294   2.885  0.00404 **
## CompanyHP        175.313     44.449   3.944  8.87e-05 ***
## CompanyLenovo     60.894     42.790   1.423  0.15519
## TypeNameNotebook -234.455     52.273  -4.485  8.60e-06 ***
## TypeNameUltrabook 203.684     65.418   3.114  0.00193 **
## GPUIntel         163.638     40.534   4.037  6.05e-05 ***
## GPUNvidia        227.598     47.445   4.797  2.00e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 345 on 654 degrees of freedom
## Multiple R-squared:  0.6441, Adjusted R-squared:  0.6386
## F-statistic: 118.3 on 10 and 654 DF,  p-value: < 2.2e-16

# out-of-sample R^2 for model2
pred_test <- predict(model2, newdata = test)
sse <- sum((test$Price - pred_test)^2)
sst <- sum((test$Price - mean(test$Price))^2)
R2_out <- 1 - sse/sst
R2_out

## [1] 0.5216836
```

Here we notice that our OOS  $R^2$  value and Multiple R-squared values have dropped. Although removing InventoryID and Weight reduces the OOS  $R^2$  of our model (from 0.5506981 to 0.5216836), InventoryID was removed for the above mentioned reasons as it was not significant and also has no managerial insight or impact on the model. However, removing Weight was a tradeoff between predictive power and interpretability given that it is highly correlated with Screen. If our emphasis was purely on predictive power, one could make a case to include it in the model.

```
# 3rd Model to compare with
model3 <- lm(Price ~ Screen + Memory + Company + TypeName + GPU, data = train)
summary(model3)
```

```
##
## Call:
## lm(formula = Price ~ Screen + Memory + Company + TypeName + GPU,
##     data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1094.55  -205.66   -29.65   159.11  1543.68
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    716.852    270.578   2.649  0.00826 **
## Screen         -34.580     15.850  -2.182  0.02949 *
## Memory          90.349      4.485  20.144 < 2e-16 ***
## CompanyDell    127.121     43.497   2.923  0.00359 **
## CompanyHP     177.134     44.660   3.966 8.11e-05 ***
## CompanyLenovo  64.252     42.981   1.495  0.13542
## TypeNameNotebook -228.566    52.483  -4.355 1.54e-05 ***
## TypeNameUltrabook 209.510    65.702   3.189  0.00150 **
## GPUIntel       162.281     40.728   3.985 7.52e-05 ***
## GPUNvidia      223.443     47.652   4.689 3.34e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 346.7 on 655 degrees of freedom
## Multiple R-squared:  0.64, Adjusted R-squared:  0.6351
## F-statistic: 129.4 on 9 and 655 DF, p-value: < 2.2e-16
```

```
# out-of-sample R2 for model3
pred_test <- predict(model3, newdata = test)
sse <- sum((test$Price - pred_test)^2)
sst <- sum((test$Price - mean(test$Price))^2)
R2_out <- 1 - sse/sst
R2_out
```

```
## [1] 0.5225497
```

Here I noticed that in the 3rd model, removing rating improves our OOS  $R^2$  but worsens our Multiple R-squared. But I still chose to include it as it remains statistically significant in our model and also provides managerial impact and is intuitively a factor that consumers consider when thinking about price.

After considering the models, I decided to go ahead with Model2 (i.e dropping InventoryID and Weight, but not dropping Rating) as this configuration provides much more interpretability at a slightly low prediction power. It's important to know the requirements of our client so we can decide the tradeoff between predictive power and explainability/managerial sense.

```
model2 <- lm(Price ~ Screen + Memory + Rating + Company + TypeName + GPU, data = train)
summary(model2)
```

```
##
## Call:
```

```
## lm(formula = Price ~ Screen + Memory + Rating + Company + TypeName +
##     GPU, data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1040.73  -194.38   -36.05   166.39  1559.87
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      823.027    272.079   3.025  0.00258 **
## Screen          -36.598     15.791  -2.318  0.02077 *
## Memory           90.206      4.464  20.208 < 2e-16 ***
## Rating          -12.815      4.710  -2.721  0.00669 **
## CompanyDell      124.905     43.294   2.885  0.00404 **
## CompanyHP       175.313     44.449   3.944 8.87e-05 ***
## CompanyLenovo    60.894     42.790   1.423  0.15519
## TypeNameNotebook -234.455     52.273  -4.485 8.60e-06 ***
## TypeNameUltrabook 203.684     65.418   3.114  0.00193 **
## GPUIntel         163.638     40.534   4.037 6.05e-05 ***
## GPUNvidia        227.598     47.445   4.797 2.00e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 345 on 654 degrees of freedom
## Multiple R-squared:  0.6441, Adjusted R-squared:  0.6386
## F-statistic: 118.3 on 10 and 654 DF,  p-value: < 2.2e-16
```

(B)

In the final model, the effects of Screen, Memory, Company, Rating, Ultrabook type, and GPU are managerially sensible, as they align with expectations that screen size, RAM, premium designs, brand value, GPUs have a strong impact on laptop prices. However, the negative effect of Screen size and Rating seems to be counterintuitive. So it is imperative to investigate those features further in depth and look carefully at the feature data provided. But the other variables included do follow expectations based on their estimates and their impact on price.

(C)

Based on the model, HP has the highest effect on laptop price adding 175.313 more than Asus holding other factors constant, commanding the largest premium over the other brands. Even though Lenovo's coefficients are the lowest among the listed coefficients, one can argue that Asus has the smallest effect, since it is the baseline and all other manufacturers have higher estimated coefficients. While Lenovo's effect is positive (60.894), it is not statistically significant, which suggests Lenovo laptops are priced similarly to Asus on average.

(D)

```
pred_test <- predict(model2, newdata = test)
sse <- sum((test$Price - pred_test)^2)
sst <- sum((test$Price - mean(test$Price))^2)
R2_out <- 1 - sse/sst
R2_out # out-of-sample R2
```

```
## [1] 0.5216836
```

Interpretation: This means that when predicting laptop prices on unseen test data, the model explains about 52.16836% of the variation in prices compared to simply predicting the mean price for all laptops.

Formally, out-of-sample  $R^2$  is defined as  $R^2 = 1 - (SSE/SST)$ , where SSE is the sum of squared prediction errors on the test set and SST is the total sum of squared deviations of the test set prices from their mean. An out-of-sample  $R^2$  of 0.5216836 indicates that the model explains 52.16836% of the variation in laptop prices in the test data, compared to a baseline model that always predicts the average price.

(E)

```
newlap <- data.frame(
  InventoryID = 950,
  Screen = 15.6,
  Memory = 6,
  Weight = 3,
  Rating = 8,
  Company = factor("Asus", levels = levels(train$Company)),
  TypeName = factor("Ultrabook", levels = levels(train$TypeName)),
  GPU = factor("Intel", levels = levels(train$GPU))
)

# Prediction with prediction interval (includes error variance)
pred <- predict(model2, newdata = newlap, interval = "prediction", level = 0.95)
pred

##          fit          lwr          upr
## 1 1058.133 371.3338 1744.931

p <- predict(model2, newdata = newlap, se.fit = TRUE)
sigma2 <- summary(model2)$sigma^2
se_pred <- sqrt(p$se.fit^2 + sigma2)
df <- model2$df.residual

t_stat <- (1100 - p$fit) / se_pred
prob_gt_1100 <- 1 - pt(t_stat, df = df)
prob_gt_1100

##          1
## 0.4523782
```

For this laptop, the model predicts an average price of 1,058.133 euros, with a 95% prediction interval ranging from 371.3338 euros to 1,744.931 euros. The probability that the actual price exceeds €1,100 is 45.23782%. This calculation is based on the assumptions that regression errors are normally distributed (so the t-distribution approximation for the probability is valid), homoscedastic, and independent, and that the model is correctly specified.

(F)

```
model_mem_gpu <- lm(Price ~ Screen + Memory + Rating + TypeName + Company +
  GPU + Memory:GPU, data = train)
summary(model_mem_gpu)

##
## Call:
## lm(formula = Price ~ Screen + Memory + Rating + TypeName + Company +
##     GPU + Memory:GPU, data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1086.71  -194.71   -30.79   158.38  1569.31
```

```
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    1044.932     277.494   3.766 0.000181 ***
## Screen         -38.107      15.765  -2.417 0.015915 *
## Memory          60.243       9.964   6.046 2.5e-09 ***
## Rating         -13.055       4.675  -2.792 0.005384 **
## TypeNameNotebook -211.650     54.205  -3.905 0.000104 ***
## TypeNameUltrabook 213.135     69.884   3.050 0.002382 **
## CompanyDell      121.588     42.982   2.829 0.004816 **
## CompanyHP        165.627     44.206   3.747 0.000195 ***
## CompanyLenovo     61.037     42.503   1.436 0.151466
## GPUIntel        -70.059     90.786  -0.772 0.440573
## GPUNvidia       -78.966    101.655  -0.777 0.437557
## Memory:GPUIntel   32.854     11.941   2.751 0.006102 **
## Memory:GPUNvidia  39.956     11.632   3.435 0.000631 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 342.4 on 652 degrees of freedom
## Multiple R-squared:  0.6505, Adjusted R-squared:  0.6441
## F-statistic: 101.1 on 12 and 652 DF,  p-value: < 2.2e-16

pred_test <- predict(model_mem_gpu, newdata = test)
sse <- sum((test$Price - pred_test)^2)
sst <- sum((test$Price - mean(test$Price))^2)
R2_out <- 1 - sse/sst
R2_out  # out-of-sample R2

## [1] 0.5180113
```

The interaction terms show that the effect of Memory depends on the GPU type installed. For AMD laptops, each extra unit of memory adds about 60.243 euros to price. For Intel and Nvidia laptops, the memory premium is much higher, about 93.097 euros and 100.199 euros per memory unit, respectively. This suggests that additional RAM is valued more when paired with stronger GPUs. Compared to the model in part (a), this specification highlights brand-technology complementarities, and shows that consumers value memory more when paired with GPUs, though overall model fit (adjusted  $R^2$ ) is slightly lower. When I added the Memory  $\times$  GPU interaction, the adjusted  $R^2$  increased slightly (0.639 to 0.644), but the out-of-sample  $R^2$  decreased marginally (0.522 to 0.518) which suggests no predictive gain, but the model provides richer interpretation: the price premium for additional RAM depends on GPU type.

(G)

```
train$Company <- factor(train$Company, levels = c("Asus", "Dell", "HP", "Lenovo"))
train$TypeName <- factor(train$TypeName, levels = c("Gaming", "Notebook", "Ultrabook"))
train$GPU <- factor(train$GPU, levels = c("AMD", "Intel", "Nvidia"))

# Model with GPU  $\times$  Company interaction
model_gpu_company <- lm(
  Price ~ Screen + Memory + Rating + TypeName + GPU + Company + GPU:Company,
  data = train
)
summary(model_gpu_company)

##
```



```
## Call:
## lm(formula = Price ~ Screen + Memory + Rating + TypeName + GPU +
##     Company + GPU:Company, data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1162.7  -201.0   -16.2   176.6  1515.5
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    1101.353     293.294   3.755 0.000189 ***
## Screen         -36.089      15.640  -2.307 0.021343 *
## Memory          88.693       4.388  20.212 < 2e-16 ***
## Rating        -12.318       4.623  -2.665 0.007901 **
## TypeNameNotebook -163.473     53.046  -3.082 0.002145 **
## TypeNameUltrabook 295.321     66.385   4.449 1.02e-05 ***
## GPUIntel       -270.926     141.486  -1.915 0.055950 .
## GPUNvidia      -56.560     135.229  -0.418 0.675900
## CompanyDell    -253.432     138.346  -1.832 0.067429 .
## CompanyHP      -186.109     143.655  -1.296 0.195600
## CompanyLenovo  -297.136     151.692  -1.959 0.050563 .
## GPUIntel:CompanyDell 394.962     151.783   2.602 0.009476 **
## GPUNvidia:CompanyDell 517.452     153.830   3.364 0.000814 ***
## GPUIntel:CompanyHP  472.646     155.538   3.039 0.002471 **
## GPUNvidia:CompanyHP 178.070     162.639   1.095 0.273978
## GPUIntel:CompanyLenovo 490.746     163.354   3.004 0.002766 **
## GPUNvidia:CompanyLenovo 242.211     162.478   1.491 0.136519
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 337.5 on 648 degrees of freedom
## Multiple R-squared:  0.6625, Adjusted R-squared:  0.6541
## F-statistic: 79.49 on 16 and 648 DF, p-value: < 2.2e-16

pred_test <- predict(model_gpu_company, newdata = test)
sse <- sum((test$Price - pred_test)^2)
sst <- sum((test$Price - mean(test$Price))^2)
R2_out <- 1 - sse/sst
R2_out  # out-of-sample R2

## [1] 0.5331728
```

The GPU \* Company interaction terms show that the price impact of GPUs varies by manufacturer. For Asus (baseline), Intel and Nvidia GPUs do not add value, but for Dell, HP, and Lenovo, Intel GPUs command strong positive premiums of about 394-491 euros. Nvidia GPUs also add large premiums at Dell but not at HP or Lenovo. Compared to the model in part (a), this specification provides richer insight into brand-specific GPU pricing and slightly improves both adjusted R<sup>2</sup> (0.639 to 0.654) and out-of-sample R<sup>2</sup> (0.522 to 0.533).